



Analysis of SVD Deep Neural Network parametrization

“Da kto takoi etot et al?!” team

Patakin Nikolay,
Yermekova Assel,
Shushkova Varvara,
Ildar Gabdrakhmanov,
Pudiyakov Yaroslav

Theoretical Part

Problem

Neural Networks



Time consuming
operations!

- 1) Matrix inversion
- 2) Matrix Determinant
- 3) Spectral normalizations

Theoretical Part

Problem



Σ - diagonal
 U, V - orthogonal

What is the problem? →

Gradient
Descent Update
of Weight Matrix

$$\Sigma' = \Sigma - \eta \nabla_{\Sigma}$$

$$U' = U - \eta \nabla_U$$

$$V' = V - \eta \nabla_V$$

Theoretical Part

Problem

Neural Networks



Time consuming
operations

- 1) Matrix inversion
- 2) Matrix Determinant
- 3) Spectral normalizations

Solution



SVD
decomposition

$$W = U\Sigma V^T$$

Σ - diagonal
 U, V - orthogonal

Theoretical Part

Problem

Neural Networks

Time consuming
operations

- 1) Matrix inversion
- 2) Matrix Determinant
- 3) Spectral normalizations

Solution

SVD
decomposition

$$W = U \Sigma V^T$$

Σ - diagonal
 U, V - orthogonal

What is the
problem?

Gradient
Descent Update
of Weight Matrix

$$\Sigma' = \Sigma - \eta \nabla_{\Sigma}$$

$$U' = U - \eta \nabla_U$$

$$V' = V - \eta \nabla_V$$

Loss of
orthogonality of
 U, V matrices

Theoretical Part

To preserve U, V orthogonality:

$$U \in \mathbb{R}^{d \times d}$$

Householder
matrices

$$U = \prod_{i=1}^d H_i$$

$$H_i = I - 2 \frac{v_i v_i^T}{\|v_i\|_2^2} \quad v_i \in \mathbb{R}^d.$$

Why Householder matrices are good?



U **remains orthogonal** under gradient descent update at i step



It allows to perform gradient descent to **preserve the SVD** of W during gradient descent updates

All products of Householder matrices are orthogonal

Any $d \times d$ orthogonal matrix can be decomposed as a product of d Householder matrices

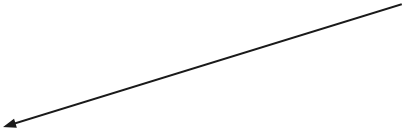
Allows to perform gradient descent over orthogonal matrices

What is the FastH algorithm?

Need to calculate:

$$UX = H_1 \cdots (H_{d-1}(H_d \cdot X))$$

$X \in \mathbb{R}^{d \times m}$

- 
- Sequential

$$O(d^2 m)$$

$O(d)$ sequential vector-vector
operations

What is the FastH algorithm?

Need to calculate:

$$UX = H_1 \cdots (H_{d-1}(H_d \cdot X))$$

$X \in \mathbb{R}^{d \times m}$

- Sequential

$$O(d^2 m)$$

$O(d)$ sequential vector-vector
operations

- FastH

$$O(d^2 m)$$

$O(\frac{d}{m} + m)$ sequential matrix
operations



FastH. Forward Pass

Input: $X \in \mathbb{R}^{d \times m}$, $d > m > 1$

H_1, \dots, H_d , $H_i \in \mathbb{R}^{d \times d}$ - Householder matrices

Want to compute: $A = H_1 \cdots H_d X$

$$H_i = I - \frac{2v_i v_i^T}{\|v_i\|_2^2}, \quad P_i = \underbrace{H_{(i-1) \cdot m + 1} \cdots H_{i \cdot m}}_{m \text{ matrices}} \quad i = 1, \dots, \frac{d}{m}$$

$$\text{Then: } A = H_1 \cdots H_d X = \underbrace{P_1 \cdots P_{\frac{d}{m}}}_{\frac{d}{m} \text{ multiplications}} X$$

d multiplications

$P_i X$ takes $O(d^2 m)$
Then compute A
takes $O(d^3)$



Still slow.... :c

Using decomposition of product $\sim O(dm^2)$

$$H_1 \cdots H_m = I - 2WY^T$$

helps us reduce complexity of computation to $O(d^2 m)$
and decrease number of matrix multiplication to $O(\frac{d}{m} + m)$



FastH. Backwards Propagation

Input: $A_1, \dots, A_{\frac{d}{m}+1}, P_1, \dots, P_{\frac{d}{m}+1}, \frac{\partial L}{\partial A_1}$ L - loss function

Want to compute: $\frac{\partial L}{\partial X}, \frac{\partial L}{\partial v_1}, \dots, \frac{\partial L}{\partial v_d}$

The backward pass of FastH has two steps:

1. Compute $\frac{\partial L}{\partial A_2}, \frac{\partial L}{\partial A_3}, \dots, \frac{\partial L}{\partial A_{\frac{d}{m}+1}}$ by $\frac{\partial L}{\partial A_{i+1}} = \left[\frac{\partial A_i}{\partial A_{i+1}} \right]^T \frac{\partial L}{\partial A_i} = P_i^T \frac{\partial L}{\partial A_i} \longrightarrow$ gradient wrt. X since $X = A_{\frac{d}{m}+1}$

2. Compute $\frac{\partial L}{\partial v_j}$ for all $j \longrightarrow$ can be split into d/m subproblems \longrightarrow can be solved in parallel \longrightarrow one subproblem for each

$$\frac{\partial L}{\partial A_i}$$


Applications

- Normalization flow models
 - Matrix determinant
 - Inverse matrix
- Spectral normalization

Normalization flow

What do we want?

simple distributions $\xrightarrow{\text{map}}$ complex distributions



How?

By a sequence of invertible and differentiable mappings

Change of variables

$$p_X(\mathbf{x}) = p_Z(f^{-1}(\mathbf{x})) \left| \det \left(\frac{\partial f^{-1}(\mathbf{x})}{\partial \mathbf{x}} \right) \right|$$

For invertible matrix:

$$p_X(\mathbf{x}) = p_Z(\mathbf{z}) \left| \det \left(\frac{\partial f(\mathbf{z})}{\partial \mathbf{z}} \right) \right|^{-1}$$

1. $f(x)$ has to be invertible
2. want fast calculation of determinant of Jacobian matrix

$$\det \left(\frac{\partial f^{-1}(x)}{\partial x} \right)$$

Flow-based Generative Models

\mathbf{x} - a high-dimensional random vector, collect a dataset D

generative process

$$\mathbf{z} \sim p_{\theta}(\mathbf{z})$$
$$\mathbf{x} = g_{\theta}(\mathbf{z})$$

f is composed of a sequence of transformations

$$f = f_1 \circ f_2 \circ \dots \circ f_K$$

$$\mathbf{x} \xleftrightarrow{f_1} \mathbf{h}_1 \xleftrightarrow{f_2} \mathbf{h}_2 \dots \xleftrightarrow{f_K} \mathbf{z}$$

normalizing flow

Normalizing flows are generative models which produce tractable distributions.


The probability density function of the model given a datapoint:

$$\log p_{\theta}(\mathbf{x}) = \log p_{\theta}(\mathbf{z}) + \log |\det(d\mathbf{z}/d\mathbf{x})|$$
$$= \log p_{\theta}(\mathbf{z}) + \sum_{i=1}^K \log |\det(d\mathbf{h}_i/d\mathbf{h}_{i-1})|$$

The log-determinant is the change in log-density when going from \mathbf{h}_{i-1} to \mathbf{h}_i under transformation f_i .

Glow: Generative Flow

Generative flow
where each step



$$(x_a, x_b) \mapsto (y_a, y_b)$$


$$y_a = x_a$$

$$y_b = s(x_a) \odot x_b + t(x_a)$$

$$\det \left(\frac{dy}{dx} \right) = \det \begin{pmatrix} I & 0 \\ \frac{dy_a}{dx_a} & \frac{dy_b}{dx_b} \end{pmatrix} = \det \left(\frac{dy_b}{dx_b} \right)$$

Glow: Generative Flow

Generative flow
where each step



$$y_{ij} = Wx_{ij}$$
$$x_{ij} \in \mathbb{R}^d, W \in \mathbb{R}^{d \times d}$$

Log-determinant

$$\det\left(\frac{dy}{dx}\right) = h \cdot w \cdot \log|\det(W)|$$

$$W = PL(U + \text{diag}(s))$$


or

$$W = U\Sigma V^T$$

Can use FastH(SVD) here

Glow: Generative Flow

Generative flow
where each step



$$(x_a, x_b) \mapsto (y_a, y_b)$$

$$y_a = x_a$$

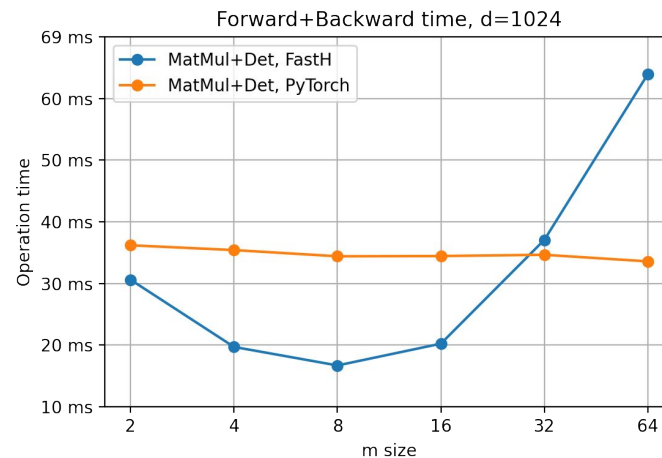

$$y_b = s(x_a) \odot x_b + t(x_a)$$

$$\forall i, j: \mathbf{y}_{i,j} = \mathbf{s} \odot \mathbf{x}_{i,j} + \mathbf{b}$$

$$\det \left(\frac{dy}{dx} \right) = \det \begin{pmatrix} I & 0 \\ \frac{dy_a}{dx_a} & \frac{dy_b}{dx_b} \end{pmatrix} = \det \left(\frac{dy_b}{dx_b} \right)$$

$$\det \left(\frac{dy}{dx} \right) = h \cdot w \cdot \text{sum}(\log |\mathbf{s}|)$$

How we can use FastH(SVD) here



How we can use FastH(SVD) here

Efficient FastH operation: $UX, U \in \mathbb{R}^{d \times d}, X \in \mathbb{R}^{d \times m}$

But! By algorithm formulation $m \ll d$

For conv 1x1 $X \in \mathbb{R}^{d \times BHW}$ B - batch size
H - height
W - width $\Rightarrow m \gg d$

Can not be applied straightforwardly!

How we can use FastH(SVD) here

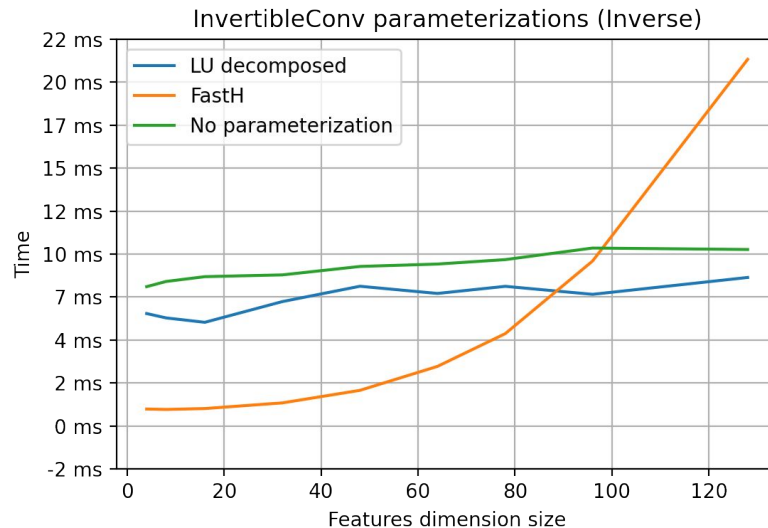

Alternative: 1) extract weight matrix:

$$W = \begin{bmatrix} U \begin{pmatrix} I_{d/2} \\ 0 \end{pmatrix} \\ U \begin{pmatrix} 0 \\ I_{d/2} \end{pmatrix} \end{bmatrix}$$

using two FastH forwards

2) Apply 2D 1x1 convolution using W matrix

How we can use FastH(SVD) here




GLOW performance

	Pytorch	FastH
Forward	160 ms	112 ms
Inverse	475 ms	113 ms
Forward+Backward	513 ms	245 ms

~7.5 min (out of 15) per epoch reduction!
(CIFAR10)

CIFAR-10

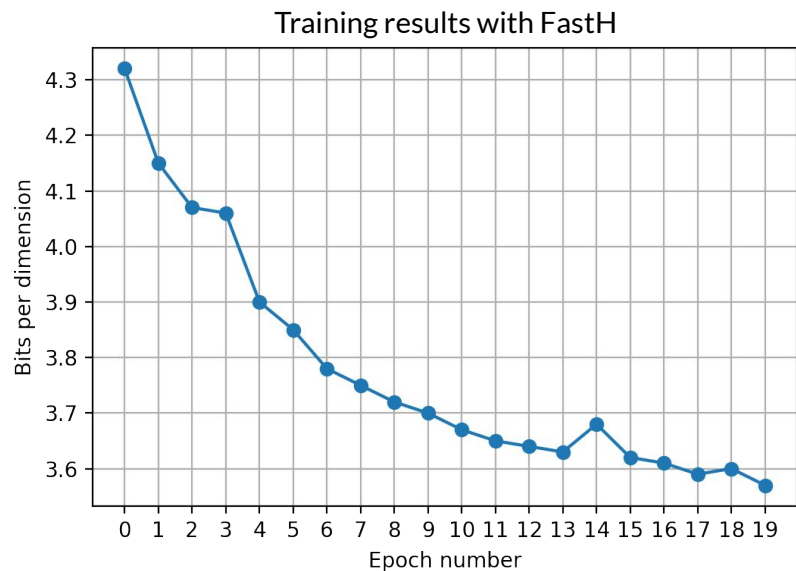
Real examples



After 20 epochs



CIFAR-10



$$BPD = \frac{NLL}{3 \ln 2 \cdot H \cdot W}$$

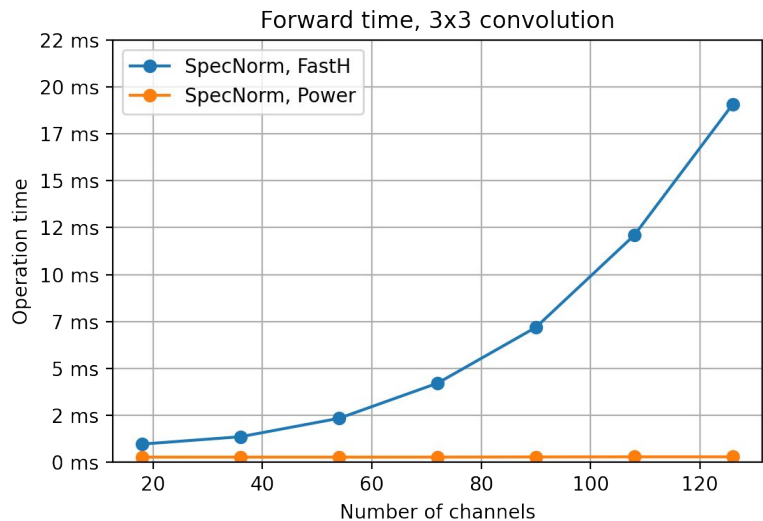
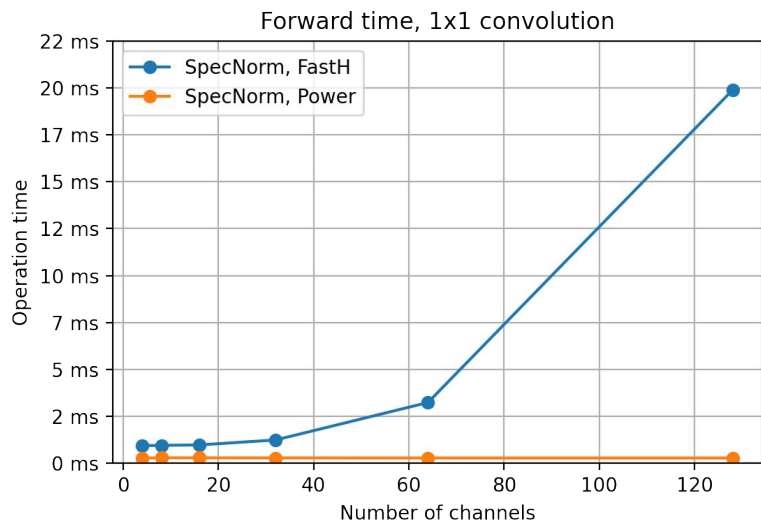
- After 20 epochs: **3.57 bpd**
- For original GLOW should be ~3.48 after 80 epochs



It works!

Spectral normalization

$$W \in \mathbb{R}^{C_{out} \times C_{in} \times K \times K} \quad \sigma(W) \rightarrow 1$$



Conclusion

- Practical speed-up can be achieved for GLOW model

BUT:

- A lot of limitations on parameters choice
- There is more powerful models (e.g. Neural Splines Flow)
- The only normalization flow model can be accelerated?

Spectral norms:

- Slower than power iteration