

Skoltech

NLA 2020 Final Project

Active Subspace of Neural Networks: Structural Analysis

Team: Passive Subspaces

Daniil Cherniavskii, Vanessa Skliarova, Aleksei Panchenko, Saveliy Galochkin



Problem statement

- **What?** - Neural network compression, using active subspace structural analysis
- **Why?** - Less memory, higher processing speed, more simple architecture
- **Hypothesis:** check that we can cut off network tail at some middle layer and approximate tail with a couple of much simpler layers with good compression rate and reasonable quality.
- **Applications** - use compressed models in applications.
- **How to measure quality?** - Compression rate - how much nnz parameters are compressed in ASNet + accuracy w/ fine tuning and w/o


[Active Subspace of Neural Networks: Structural Analysis and Universal Attacks, Chunfeng Cui, Kaiqi Zhang, Talgat Daulbaev, Julia Gusak, Ivan Oseledets, Zheng Zhang](#)

Neural Network Compression Methods

- Weight Sharing
- Network Pruning
- Low Rank Matrix & Tensor Decompositions
- Knowledge Distillation
- Quantization

Neural Network Compression Techniques

- Weight Sharing
- Network Pruning
- **Low Rank Matrix & Tensor Decompositions**
- Knowledge Distillation
- Quantization



An Survey of Neural Network Compression <https://arxiv.org/pdf/2006.03669.pdf>

Major Points

- Activate Subspace
- Structural analysis and compression of deep neural networks
- Active subspace network (ASNet)
 - The active subspace layer
 - Polynomial chaos expansion layer
- The training procedure of the ASNet
- Experimental Results
- Conclusion

Active Subspace

$c(x)$ - continuous function

$$C = \mathbb{E}[\nabla c(x) \nabla c(x)^T]$$

$$C = V \Lambda V^T$$

$$\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n), \lambda_1 \geq \dots \geq \lambda_n \geq 0$$

$$V = [V_1, V_2], \text{ where } V_1 \in \mathbb{R}^{n \times r} \text{ and } V_2 \in \mathbb{R}^{n \times (n-r)}$$

Active Subspace

$c(x)$ - continuous function

$$C = \mathbb{E}[\nabla c(x) \nabla c(x)^T]$$

$$C = V \Lambda V^T$$

$$\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n), \lambda_1 \geq \dots \geq \lambda_n \geq 0$$

$$V = [V_1, V_2], \text{ where } V_1 \in \mathbb{R}^{n \times r} \text{ and } V_2 \in \mathbb{R}^{n \times (n-r)}$$



Active Subspace

Lemma 2.2 (see [10]). Suppose $c(\mathbf{x})$ is a continuous function and \mathbf{C} is obtained from (2.1). For the matrices \mathbf{V}_1 and \mathbf{V}_2 generated by (2.3), and the reduced vector

$$(2.5) \quad \mathbf{z} = \mathbf{V}_1^T \mathbf{x} \text{ and } \tilde{\mathbf{z}} = \mathbf{V}_2^T \mathbf{x},$$

it holds that

$$(2.6) \quad \begin{aligned} \mathbb{E}_{\mathbf{x}}[\nabla_{\mathbf{z}} c(\mathbf{x})^T \nabla_{\mathbf{z}} c(\mathbf{x})] &= \lambda_1 + \cdots + \lambda_r, \\ \mathbb{E}_{\mathbf{x}}[\nabla_{\tilde{\mathbf{z}}} c(\mathbf{x})^T \nabla_{\tilde{\mathbf{z}}} c(\mathbf{x})] &= \lambda_{r+1} + \cdots + \lambda_n. \end{aligned}$$

Active Subspace

2.1. Response surface. For a fixed \mathbf{z} , the best guess for g is the conditional expectation of c given \mathbf{z} , i.e.,

$$(2.7) \quad g(\mathbf{z}) = \mathbb{E}_{\tilde{\mathbf{z}}}[c(\mathbf{x})|\mathbf{z}] = \int c(\mathbf{V}_1\mathbf{z} + \mathbf{V}_2\tilde{\mathbf{z}})\rho(\tilde{\mathbf{z}}|\mathbf{z})d\tilde{\mathbf{z}}.$$

Based on the Poincaré inequality, the following approximation error bound is obtained [10].

Lemma 2.3. *Assume that $c(\mathbf{x})$ is absolutely continuous and square integrable with respect to the probability density function $\rho(\mathbf{x})$; then the approximation function $g(\mathbf{z})$ in (2.7) satisfies*

$$(2.8) \quad \mathbb{E}[(c(\mathbf{x}) - g(\mathbf{z}))^2] \leq O(\lambda_{r+1} + \cdots + \lambda_n).$$

Structural analysis and compression of deep neural networks


$$f(x_0) = f_L(f_{L-1} \dots (f_1(x_0))) \quad f(x_0) = f_{post}^l(f_{pre}^l(x_0))$$

Definition 3.1. Suppose $\mathbf{\Lambda}$ is computed by (2.2). For any layer index $1 \leq l \leq L$, we define the number of active neurons $n_{l,AS}$ as follows:


$$(3.4) \quad n_{l,AS} = \arg \min \left\{ i : \frac{\lambda_1 + \dots + \lambda_i}{\lambda_1 + \dots + \lambda_{n_l}} \geq 1 - \epsilon \right\},$$

where $\epsilon > 0$ is a user-defined threshold.

Active subspace network (ASNet)



(a) A deep neural network



(b) The proposed ASNet

Figure 2. (a) *The original deep neural network.* (b) *The proposed ASNet with three parts: a pre-model, an active subspace (AS) layer, and a polynomial chaos expansion (PCE) layer.*

The active subspace layer

Algorithm 3.2. The frequent direction algorithm for computing the active subspace.

Input: A dataset with m_{AS} input samples $\{\mathbf{x}_0^j\}_{j=1}^{m_{AS}}$, a pre-model $f_{\text{pre}}^l(\cdot)$, a subroutine for computing $\nabla c_l(\mathbf{x})$, and the dimension of truncated singular value decomposition r .

- 1: Select r samples \mathbf{x}_0^i , compute $\mathbf{x}^i = f_{\text{pre}}^l(\mathbf{x}_0^i)$, and construct an initial matrix $\mathbf{S} \leftarrow [\nabla c_l(\mathbf{x}^1), \dots, \nabla c_l(\mathbf{x}^r)]$.
- 2: **for** $t=1, 2, \dots$, **do**
- 3: Compute the singular value decomposition $\mathbf{V}\mathbf{\Sigma}\mathbf{U}^T \leftarrow \text{svd}(\mathbf{S})$, where $\mathbf{\Sigma} = \text{diag}(\sigma_1, \dots, \sigma_r)$.
- 4: If the maximal number of samples m_{AS} is reached, stop.
- 5: Update \mathbf{S} by the soft-thresholding (3.8). $\rightarrow \mathbf{S} \leftarrow \mathbf{V}\sqrt{\mathbf{\Sigma}^2 - \sigma_r^2}$
- 6: Get a new sample $\mathbf{x}_0^{\text{new}}$, compute $\mathbf{x}^{\text{new}} = f_{\text{pre}}^l(\mathbf{x}_0^{\text{new}})$, and replace the last column of \mathbf{S} (now all zeros) by the gradient vector $\mathbf{S}(:, r) \leftarrow \nabla c_l(\mathbf{x}^{\text{new}})$.
- 7: **end for**

Output: The projection matrix $\mathbf{V} \in \mathbb{R}^{n_l \times r}$ and the singular values $\mathbf{\Sigma} \in \mathbb{R}^{r \times r}$.


$$\hat{n}_{l,AS} = \arg \min \left\{ i : \frac{\sqrt{\sigma_1^2 + \dots + \sigma_i^2}}{\sqrt{\sigma_1^2 + \dots + \sigma_r^2}} \geq 1 - \epsilon \right\}.$$

Polynomial Chaos Expansion Layer



$$\hat{\mathbf{y}} \approx \sum_{|\alpha|=0}^p \mathbf{c}_{\alpha} \phi_{\alpha}(\mathbf{z}), \text{ where } |\alpha| = \alpha_1 + \dots + \alpha_d.$$

$$\min_{\{\mathbf{c}_{\alpha}\}} \frac{1}{m_{\text{PCE}}} \sum_{j=1}^{m_{\text{PCE}}} \left\| \mathbf{y}^j - \sum_{|\alpha|=0}^p \mathbf{c}_{\alpha} \phi_{\alpha}(\mathbf{z}^j) \right\|^2.$$



The training procedure of the ASNet

Algorithm 3.1. The training procedure of the ASNet.

Input: A pretrained deep neural network, the layer index l , and the number of active neurons r .

Step 1 **Initialize the active subspace layer.** The active subspace layer is a linear projection where the projection matrix $\mathbf{V}_1 \in \mathbb{R}^{n \times r}$ is computed by Algorithm 3.2.

If r is not given, we use $r = n_{\text{AS}}$ defined in (3.4) by default.

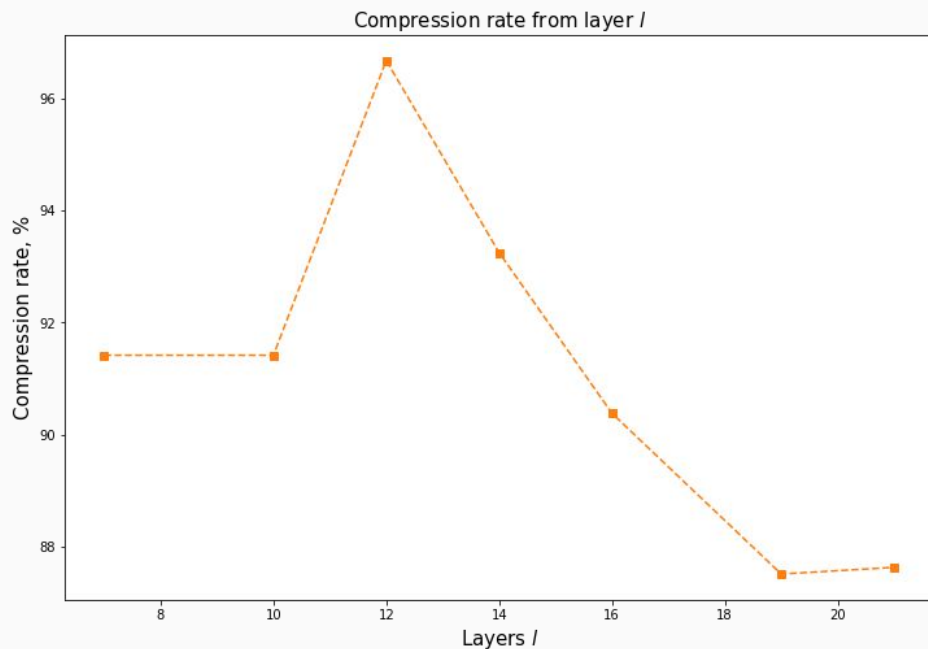
Step 2 **Initialize the polynomial chaos expansion layer.** The polynomial chaos expansion layer is a nonlinear mapping from the reduced active subspace to the outputs, as shown in (3.10). The weights \mathbf{c}_α is computed by (3.12).

Step 3 **Construct the ASNet.** Combine the pre-model (the first l layers of the deep neural network) with the active subspace and polynomial chaos expansion layers as a new network, referred to as ASNet.

Step 4 **Fine-tuning.** Retrain all the parameters in pre-model, active subspace layer, and polynomial chaos expansion layer in ASNet for several epochs by a proper (variant of) the stochastic gradient descent algorithm.


Output: A new network ASNet

Experimental Results: VGG19




Experimental Results: VGG19

Eigenvalues values for different l

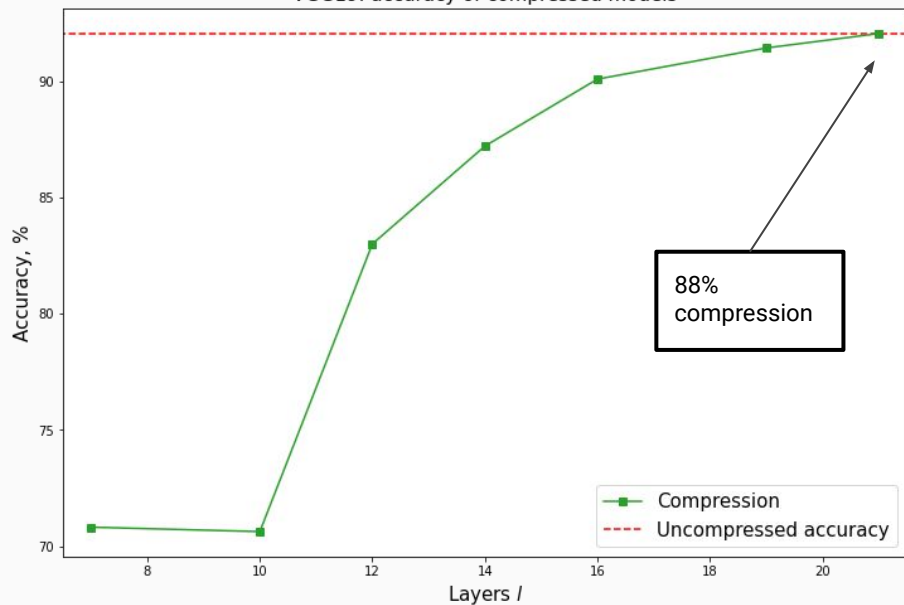


Dimension of active subspace for different l




Experimental Results: VGG19

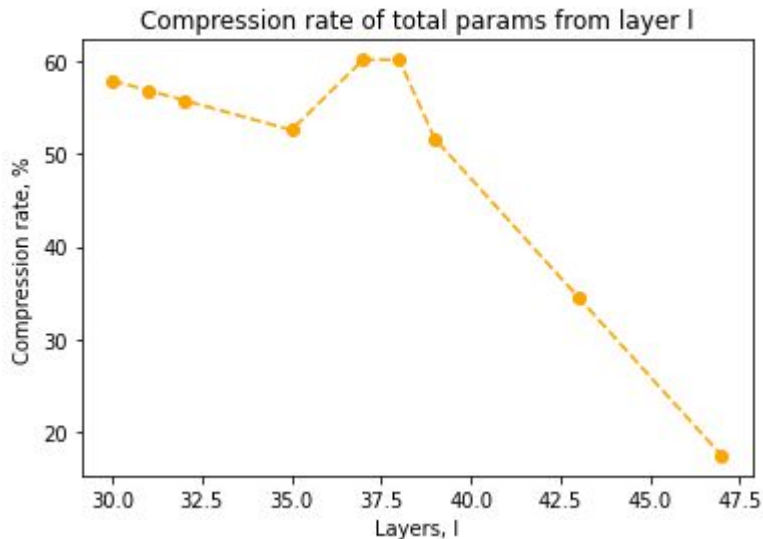
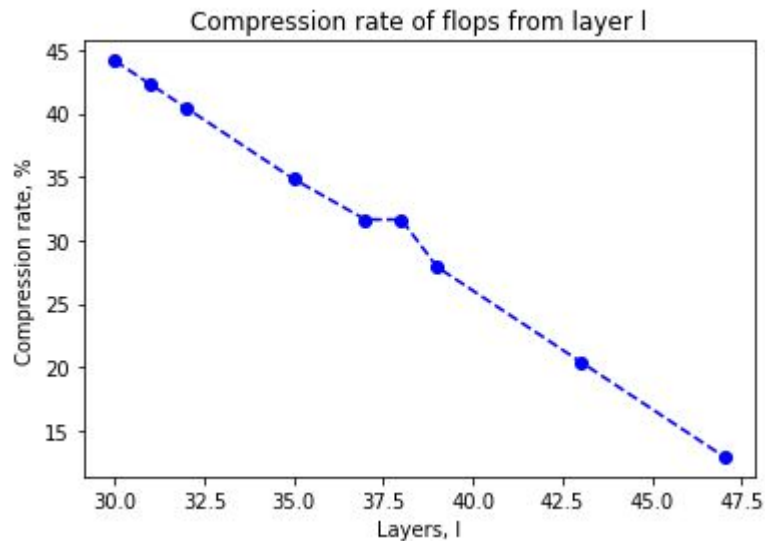
VGG19: accuracy of compressed models



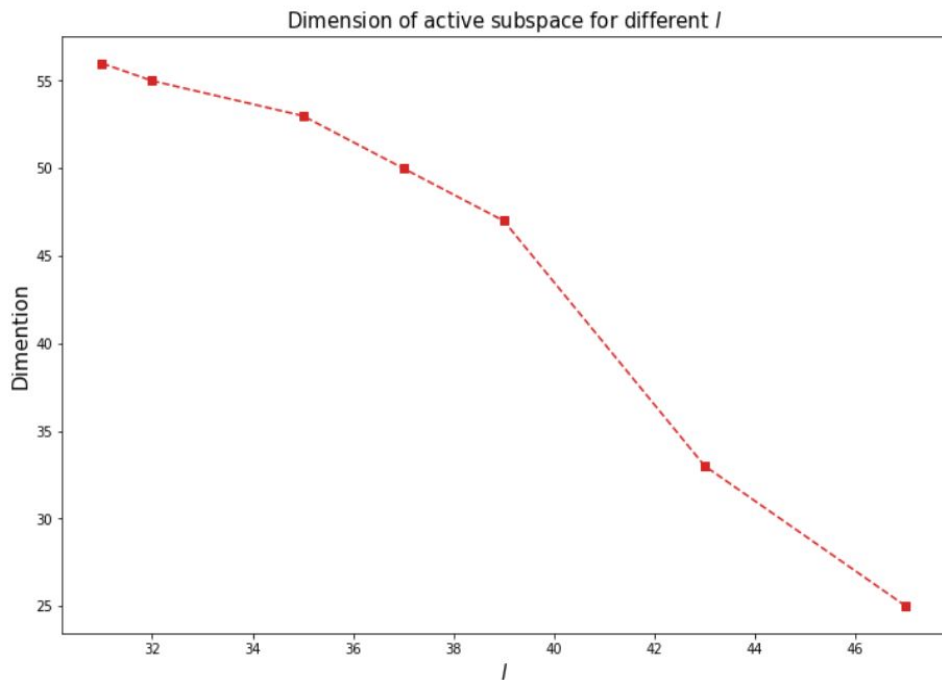
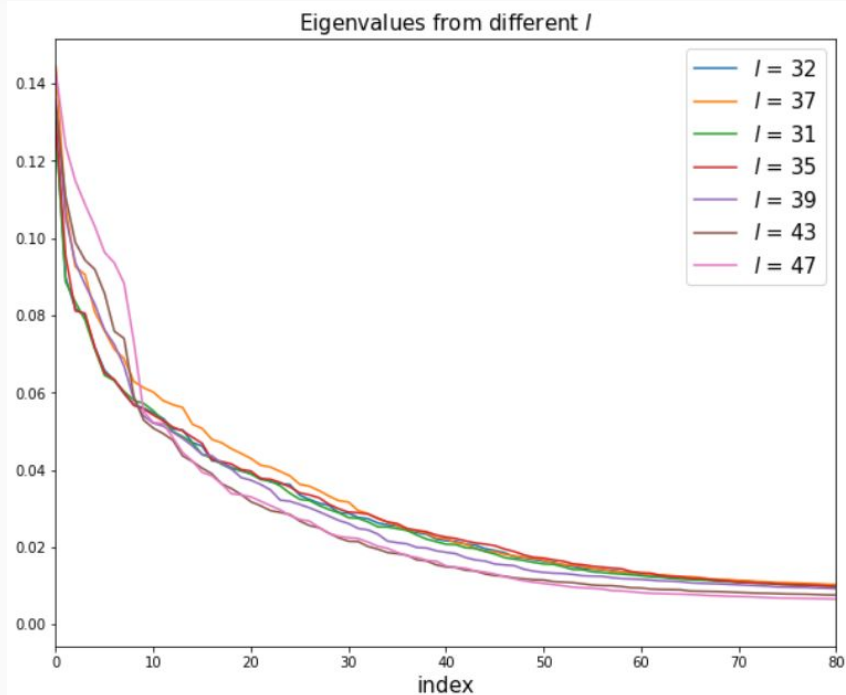
VGG19: accuracy of compressed finetuned models



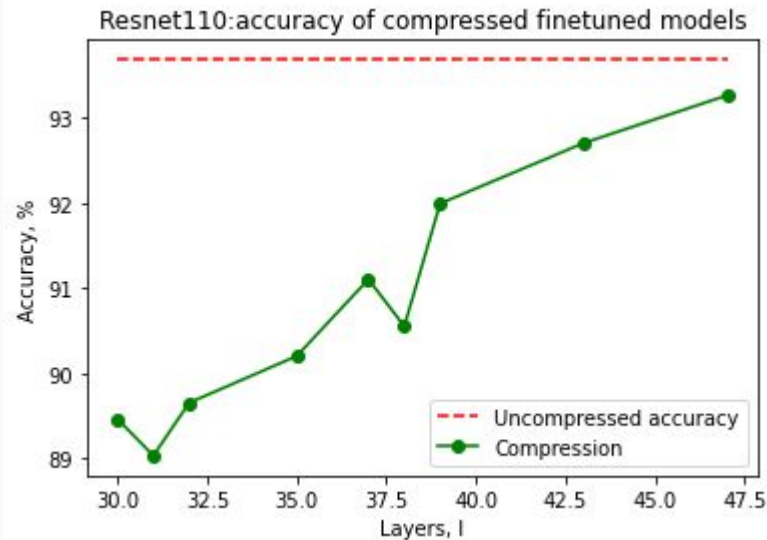

Experimental Results: Resnet110



Experimental Results: Resnet110



Experimental Results: Resnet110



Summary

- Compression approach from [original paper](#) was reproduced for **VGG19** and **Resnet110** architectures
- Compression worked differently for different network topologies, cut-off layer selection is important
- Minor code errors from the original repository were fixed

References

- [Active Subspace of Neural Networks: Structural Analysis and Universal Attacks, Chunfeng Cui, Kaiqi Zhang, Talgat Daulbaev, Julia Gusak, Ivan Oseledets, Zheng Zhang](#)
- ResNet repo: https://github.com/Kaktusava/NLA_ResNet

Questions?